

Module 3.0 - Real Neural Networks

Review: Chain Rule

$f(G(x))$

- $z_1 = g^1(x), z_2 = g^2(x)$
- $d_1 = f'_{z_1}(z_1, z_2), d_2 = f'_{z_2}(z_1, z_2)$
- $f'_x(G(x)) = d_1 g_x'^1(x) + d_2 g_x'^2(x)$

x

Processing math: 100%



Review: Chain Rule

$$f(G(x))$$

- $z_1 = g^1(x), z_2 = g^2(x), \dots$
- $d_1 = f'_{z_1}(z), d_2 = f'_{z_2}(z), \dots$
- $f'_x(G(x)) = \sum_i d_i g_x'^i(x)$

Tensor Functions

Think of it as many functions with many arguments

$$G(x) = [G^1(x_1, \dots), G^2(x_1, \dots), \dots, G^N(x_1, \dots)]$$

Derivative

Derivative of i 'th output wrt j 'th input

$$G_{x_j}^{'i}(x)$$

Full Chain Rule For Gradients

$$f(G(x))$$

- $z_1 = G^1(x), z_2 = G^2(x), \dots$
- $d_1 = f'_{z_1}(z), d_2 = f'_{z_2}(z), \dots$
- $f'_{x_j}(G(x)) = \sum_i d_i G'^i_{x_j}(x)$

Backward Function

Backward function needs to compute:

- d_i - tensor
- $G_{x_j}^{'i}$ change in i

$$\sum_i d_i G_{x_j}^{'i}(x)$$

Special Function: Map

- $G_{x_j}'^i(x) = 0$ if $i \neq j$
- $f_{x_j}'(G(x)) = d_i g_{x_j}'^j(x)$

Implies:

- $f_{x_i}'(G(x)) = d_i G_{x_i}'^i(x)$

Map Gradient

Example: Tensor Negation

- $G^i(x) = -x_i$
- $G_{x_i}^{'i}(x) = -1$
- $f_{x_i}'(G(x)) = -d_i$

Example: Tensor Negation

```
class Neg(minitorch.Function):  
    @staticmethod  
    def forward(ctx, t1: Tensor) -> Tensor:  
        return t1.f.neg_map(t1)  
  
    @staticmethod  
    def backward(ctx, d: Tensor) -> Tensor:  
        return d.f.neg_map(d)
```


Example: Tensor Inversion


- $G^i(x) = 1/x_i$
- $G_{x_i}^{'i}(x) = -(x_i)^{-2}$
- $f_{x_i}'(G(x)) = -(x_i)^{-2} * d_i$

Example: Inv

```
class Inv(minitorch.Function):
    @staticmethod
    def forward(ctx, t1: Tensor) -> Tensor:
        ctx.save_for_backward(t1)
        return t1.f.inv_map(t1)

    @staticmethod
    def backward(ctx, d: Tensor) -> Tensor:
        (t1,) = ctx.saved_values
        return d.f.inv_back_zip(t1, d)
```


Zip Gradient



Example: Tensor Inversion

- $G^i(x, y) = x_i + y_i$
- $G_{x_i}^{'i}(x, y) = 1$
- $f_{x_i}'(G(x)) = d_i$

Example: Add


```
class Add(minitorch.Function):
    @staticmethod
    def forward(ctx, t1: Tensor, t2: Tensor) -> Tensor:
        return t1.f.add_zip(t1, t2)

    @staticmethod
    def backward(ctx, grad_output: Tensor) -> Tuple[Tensor, Tensor]:
        return grad_output, grad_output
```


Example: Tensor Inversion

- $G(x) = \sum_i x_i$
- $G'_{x_i}(x) = 1$
- $f'_{x_i}(G(x)) = d$

Reduce Gradient



Quiz

Outline

- Training
- Simple NLP

Training

Parameter Fitting

1. Compute the loss function, :math: L(w_1, w_2, b)
2. See how small changes would change the loss
3. Update to parameters to locally reduce the loss

Batching

input

features

batch

Loss

1) Compute Loss ::

```
out = model.forward(X).view(data.N)
loss = -((out * y) + (out - 1.0) * (y - 1.0)).log()
```


Model: Math

$$\text{lin}(x; w, b) = x_1 \times w_1 + x_2 \times w_2 + b$$

$$h_1 = \text{ReLU}(\text{lin}(x; w^0, b^0))$$

$$h_2 = \text{ReLU}(\text{lin}(x; w^1, b^1))$$

$$m(x_1, x_2) = \text{lin}(h; w, b)$$

Model: Code

1) Model

```
class Network(minitorch.Module):  
    def __init__(self):  
        ...  
        self.layer1 = Linear(2, HIDDEN)  
        self.layer2 = Linear(HIDDEN, HIDDEN)  
        self.layer3 = Linear(HIDDEN, 1)
```


Layer 1: Weight

weights

hiddens

features

Layer 1: Bias

bias

hiddens



Key Task

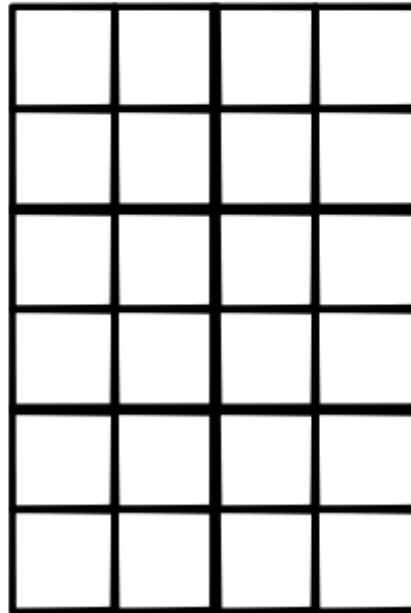
- Use broadcasting to implement the linear function
- Hint: Align `batch` x `features` x `hidden` to make it work

Layer 2: Weights

weights

hiddens

features



Compute Derivatives

Step 2

```
(loss.sum().view(1)).backward()  
print(model.layer1.w_1.value.grad)
```

weights

hiddens

features

Layer 1: Weight Grad

weights

hiddens

features

weights

hiddens

features

Update Parameters

Step 3

```
for p in model.parameters():  
    if p.value.grad is not None:  
        p.update(p.value - RATE * (p.value.grad / float(data.N)))
```


Broadcasting

- Batches
- Loss Computation
- Linear computation
- Autodifferentiation
- Gradient updates

Observations

- Exactly the same function as Module-1
- No loops within tensors

